

# Regularization I

Parameter regularization

# REGULARIZATION

**Topics:** Why we regularize.

- How do we ensure our model will perform well not just on the training data, but also on new inputs?
  - i.e. How do we ensure **generalization**?
- Our main strategy: **Regularization**
- Definition: Putting extra constraints on a machine learning model, to improved performance on the **test set** (not training set), either by encoding prior knowledge into the model, or by forcing the model to consider alternative hypotheses that explain the training data.

# REGULARIZATION


**Topics:** Why we regularize.

- Regularizers trade increased bias for reduced variance.
- An effective regularizer is one that makes a profitable trade, that is, it reduces variance significantly while not overly increasing the bias.

# REGULARIZATION

## Topics: bias-variance trade-off

- Variance of trained model: does it vary a lot if the training set changes
- Bias of trained model: is the average model close to the true solution
- Generalization error can be seen as the sum of the (squared) bias and the variance (the mean squared error - MSE)





# REGULARIZATION

In the context of deep learning:

- An overly complex model family does not necessarily include (or even come close to) the target function or the true data generating process.
- Often working with data such as images, audio sequences and text:
  - ▶ we can safely assume that the model family we are training does not include the data generating process.
- Consequence: complexity of the model is not about finding the model of the right size. i.e. the right number of parameters.
- Instead, the best fitting model is one that possesses a large number of parameters that are not entirely free to span their domain.

# REGULARIZATION

Recall: **Structured risk** (loss function we are minimizing)

$$\tilde{J}(\boldsymbol{\theta}; \mathbf{X}, \mathbf{y}) = J(\boldsymbol{\theta}; \mathbf{X}, \mathbf{y}) + \alpha \Omega(\boldsymbol{\theta})$$

Empirical Risk

Regularization Term

- $\alpha$  is a hyperparameter that balances the relative effect of the regularization term.
- For neural network models, we typically have  $\boldsymbol{\theta} = [\mathbf{w}^\top, b]^\top$
- Typical regularizers:  $\Omega(\boldsymbol{\theta}) = \frac{1}{2} \|\mathbf{w}\|_2^2$  or  $\Omega(\boldsymbol{\theta}) = \|\mathbf{w}\|_1$

# REGULARIZATION

**Topics:** L2 regularization  $\Omega(\boldsymbol{\theta}) = \frac{1}{2} ||\boldsymbol{w}||_2^2$

$$\tilde{J}(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) = J(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) + \alpha \Omega(\boldsymbol{\theta})$$

Taking gradient

$$\nabla_{\boldsymbol{w}} \tilde{J}(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y}) = \alpha \boldsymbol{w} + \nabla_{\boldsymbol{w}} J(\boldsymbol{X}, \boldsymbol{y}; \boldsymbol{w})$$

We are going to gain some insight into how this works!



# REGULARIZATION

**Topics:** L2 regularization  $\Omega(\boldsymbol{\theta}) = \frac{1}{2} ||\boldsymbol{w}||_2^2$

$$\tilde{J}(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) = J(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) + \alpha \Omega(\boldsymbol{\theta})$$

- Consider a quadratic approximation to the loss function in the neighbourhood of the empirically optimal value of the weights  $\boldsymbol{w}^*$

$$\hat{J}(\boldsymbol{\theta}) = J(\boldsymbol{w}^*) + \frac{1}{2} (\boldsymbol{w} - \boldsymbol{w}^*)^\top \boldsymbol{H} (\boldsymbol{w} - \boldsymbol{w}^*)$$

Taking gradient

$$\nabla_{\boldsymbol{w}} \hat{J}(\boldsymbol{w}) = \boldsymbol{H} (\boldsymbol{w} - \boldsymbol{w}^*)$$



# REGULARIZATION

**Topics:** L2 regularization  $\Omega(\boldsymbol{\theta}) = \frac{1}{2} ||\boldsymbol{w}||_2^2$


$$\nabla_{\boldsymbol{w}} \hat{J}(\boldsymbol{w}) = \boldsymbol{H}(\boldsymbol{w} - \boldsymbol{w}^*)$$

- Consider the effect of L2 regularization around the unregularized optimum  $\boldsymbol{w}^*$

$$\alpha \boldsymbol{w} + \boldsymbol{H}(\boldsymbol{w} - \boldsymbol{w}^*) = 0$$

$$(\boldsymbol{H} + \alpha \boldsymbol{I}) \boldsymbol{w} = \boldsymbol{H} \boldsymbol{w}^*$$

$$\tilde{\boldsymbol{w}} = (\boldsymbol{H} + \alpha \boldsymbol{I})^{-1} \boldsymbol{H} \boldsymbol{w}^*$$



# REGULARIZATION


**Topics:** L2 regularization  $\Omega(\boldsymbol{\theta}) = \frac{1}{2} ||\boldsymbol{w}||_2^2$

$$\tilde{\boldsymbol{w}} = (\boldsymbol{H} + \alpha \boldsymbol{I})^{-1} \boldsymbol{H} \boldsymbol{w}^*$$

- The presence of the regularization term moves the optimum from  $\boldsymbol{w}^*$  to  $\tilde{\boldsymbol{w}}$

- $\boldsymbol{H}$  is real and symmetric, so we can decompose it into a diagonal matrix  $\boldsymbol{\Lambda}$  and an orthogonal basis of eigenvectors,  $\boldsymbol{Q}$ , such that:

$$\boldsymbol{H} = \boldsymbol{Q} \boldsymbol{\Lambda} \boldsymbol{Q}^\top$$




# REGULARIZATION

**Topics:** L2 regularization  $\Omega(\boldsymbol{\theta}) = \frac{1}{2} ||\boldsymbol{w}||_2^2$

$$\tilde{\boldsymbol{w}} = (\boldsymbol{H} + \alpha \boldsymbol{I})^{-1} \boldsymbol{H} \boldsymbol{w}^*$$

- In the above, if we swap in  $\boldsymbol{H} = \boldsymbol{Q} \boldsymbol{\Lambda} \boldsymbol{Q}^\top$  we get:

$$\boldsymbol{Q}^\top \tilde{\boldsymbol{w}} = (\boldsymbol{\Lambda} + \alpha \boldsymbol{I})^{-1} \boldsymbol{\Lambda} \boldsymbol{Q}^\top \boldsymbol{w}^*$$






# REGULARIZATION

**Topics:** L2 regularization  $\Omega(\boldsymbol{\theta}) = \frac{1}{2} ||\boldsymbol{w}||_2^2$

$$\boldsymbol{Q}^\top \tilde{\boldsymbol{w}} = (\boldsymbol{\Lambda} + \alpha \boldsymbol{I})^{-1} \boldsymbol{\Lambda} \boldsymbol{Q}^\top \boldsymbol{w}^*$$

- The different components of  $\boldsymbol{w}^*$  are rescaled by the regularization.
- The component aligned with eigenvector  $i$  is rescaled by a factor

$$\frac{\lambda_i}{\lambda_i + \alpha}$$



# REGULARIZATION

**Topics:** L1 regularization  $\Omega(\boldsymbol{\theta}) = ||\boldsymbol{w}||_1$

$$\tilde{J}(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) = J(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) + \beta\Omega(\boldsymbol{\theta})$$

Taking gradient



$$\nabla_{\boldsymbol{w}} \tilde{J}(\boldsymbol{w}; \boldsymbol{X}, \boldsymbol{y}) = \nabla_{\boldsymbol{w}} J(\boldsymbol{X}, \boldsymbol{y}; \boldsymbol{w}) + \beta \text{sign}(\boldsymbol{w})$$

# REGULARIZATION


**Topics:** L1 regularization  $\Omega(\boldsymbol{\theta}) = ||\boldsymbol{w}||_1$

$$\tilde{J}(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) = J(\boldsymbol{\theta}; \boldsymbol{X}, \boldsymbol{y}) + \beta \Omega(\boldsymbol{\theta})$$

- Consider a quadratic approximation to the loss function in the neighbourhood of the empirically optimal value of the weights  $\boldsymbol{w}^*$

$$\hat{J}(\boldsymbol{\theta}) = J(\boldsymbol{w}^*) + \frac{1}{2}(\boldsymbol{w} - \boldsymbol{w}^*)^\top \boldsymbol{H}(\boldsymbol{w} - \boldsymbol{w}^*)$$

Taking gradient


$$\nabla_{\boldsymbol{w}} \hat{J}(\boldsymbol{w}) = \boldsymbol{H}(\boldsymbol{w} - \boldsymbol{w}^*)$$

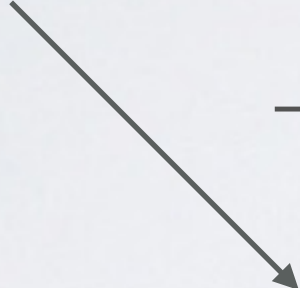


# REGULARIZATION

**Topics:** L1 regularization  $\Omega(\boldsymbol{\theta}) = ||\boldsymbol{w}||_1$

$$\hat{J}(\boldsymbol{\theta}) = J(\boldsymbol{w}^*) + \frac{1}{2}(\boldsymbol{w} - \boldsymbol{w}^*)^\top \boldsymbol{H}(\boldsymbol{w} - \boldsymbol{w}^*)$$

Taking gradient


$$\nabla_{\boldsymbol{w}} \hat{J}(\boldsymbol{w}) = \boldsymbol{H}(\boldsymbol{w} - \boldsymbol{w}^*)$$

- We will also make the further simplifying assumption that the Hessian is diagonal,  $\boldsymbol{H} = \text{diag}([\gamma_1, \dots, \gamma_N])$ , where each  $\gamma_i > 0$

# REGULARIZATION

**Topics:** L1 regularization  $\Omega(\boldsymbol{\theta}) = ||\boldsymbol{w}||_1$

- Under these assumptions the objective simplifies to a system of equations:

$$\tilde{J}(\boldsymbol{w}_i; \boldsymbol{X}, \boldsymbol{y}) = \frac{1}{2}\gamma_i(\boldsymbol{w}_i - \boldsymbol{w}_i^*)^2 + \beta|\boldsymbol{w}_i|.$$

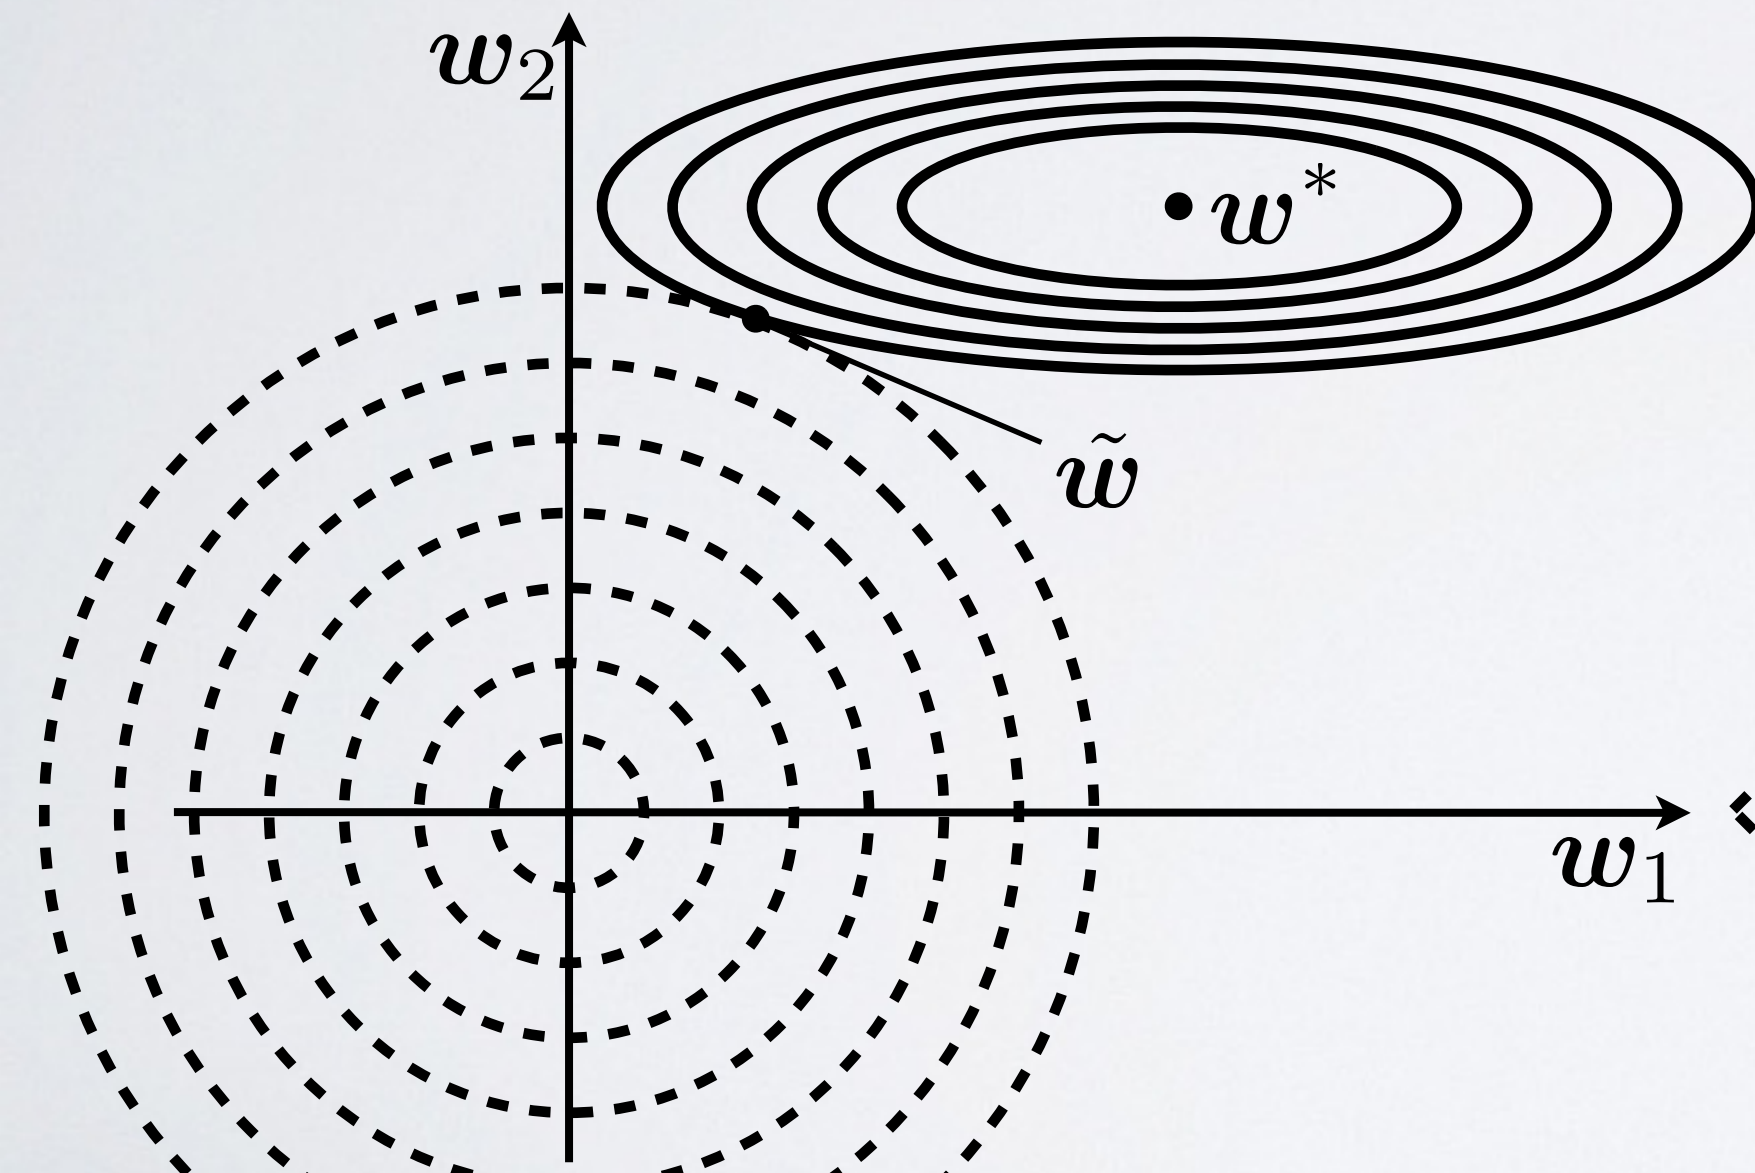
- Which admits an optimal solution (for each dimension) in the following form:

$$\boldsymbol{w}_i = \text{sign}(\boldsymbol{w}_i^*) \max(|\boldsymbol{w}_i^*| - \frac{\beta}{\gamma_i}, 0)$$

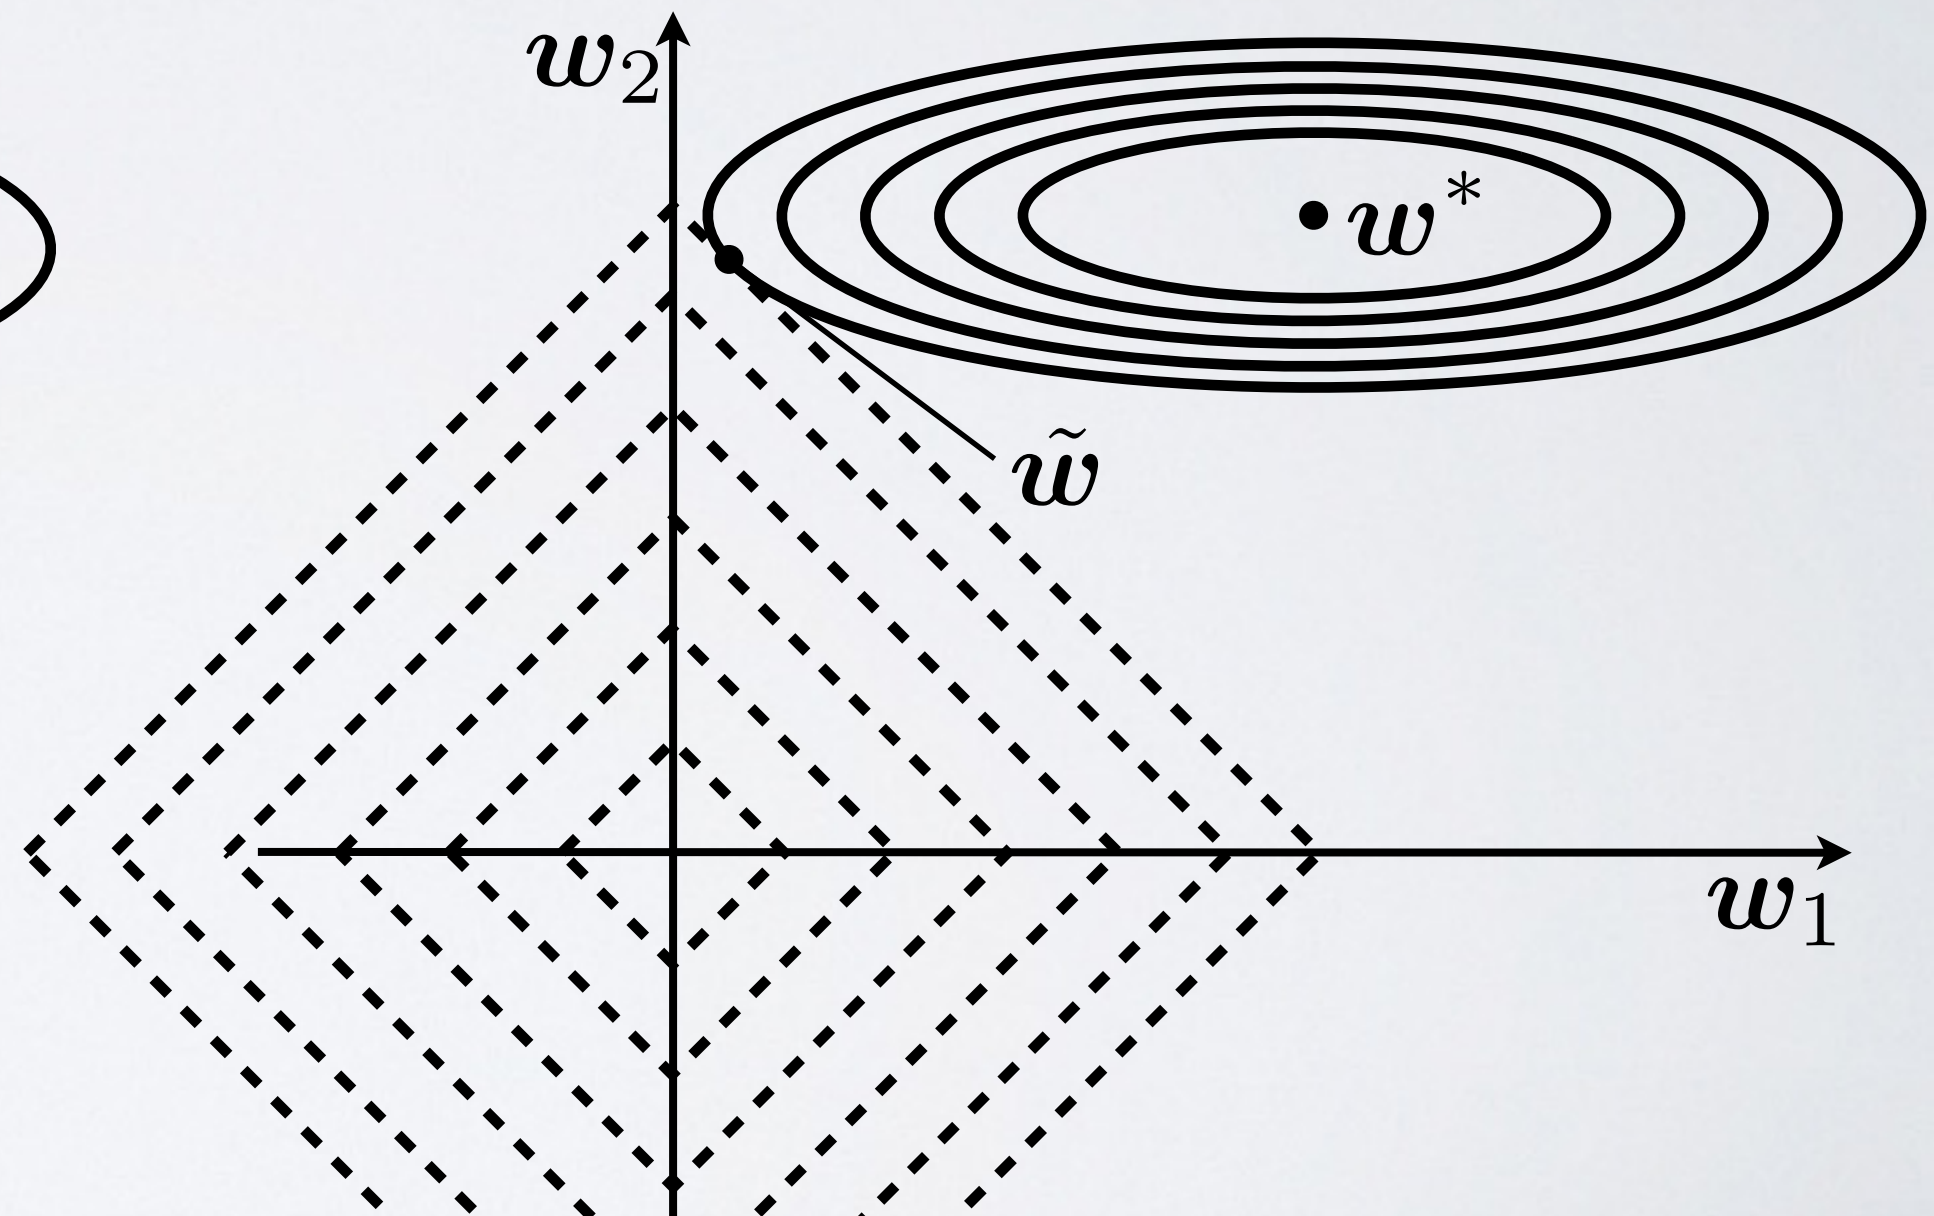
# REGULARIZATION

**Topics:** L1 regularization  $\Omega(\theta) = ||\mathbf{w}||_1$

L2 regularization



L1 regularization

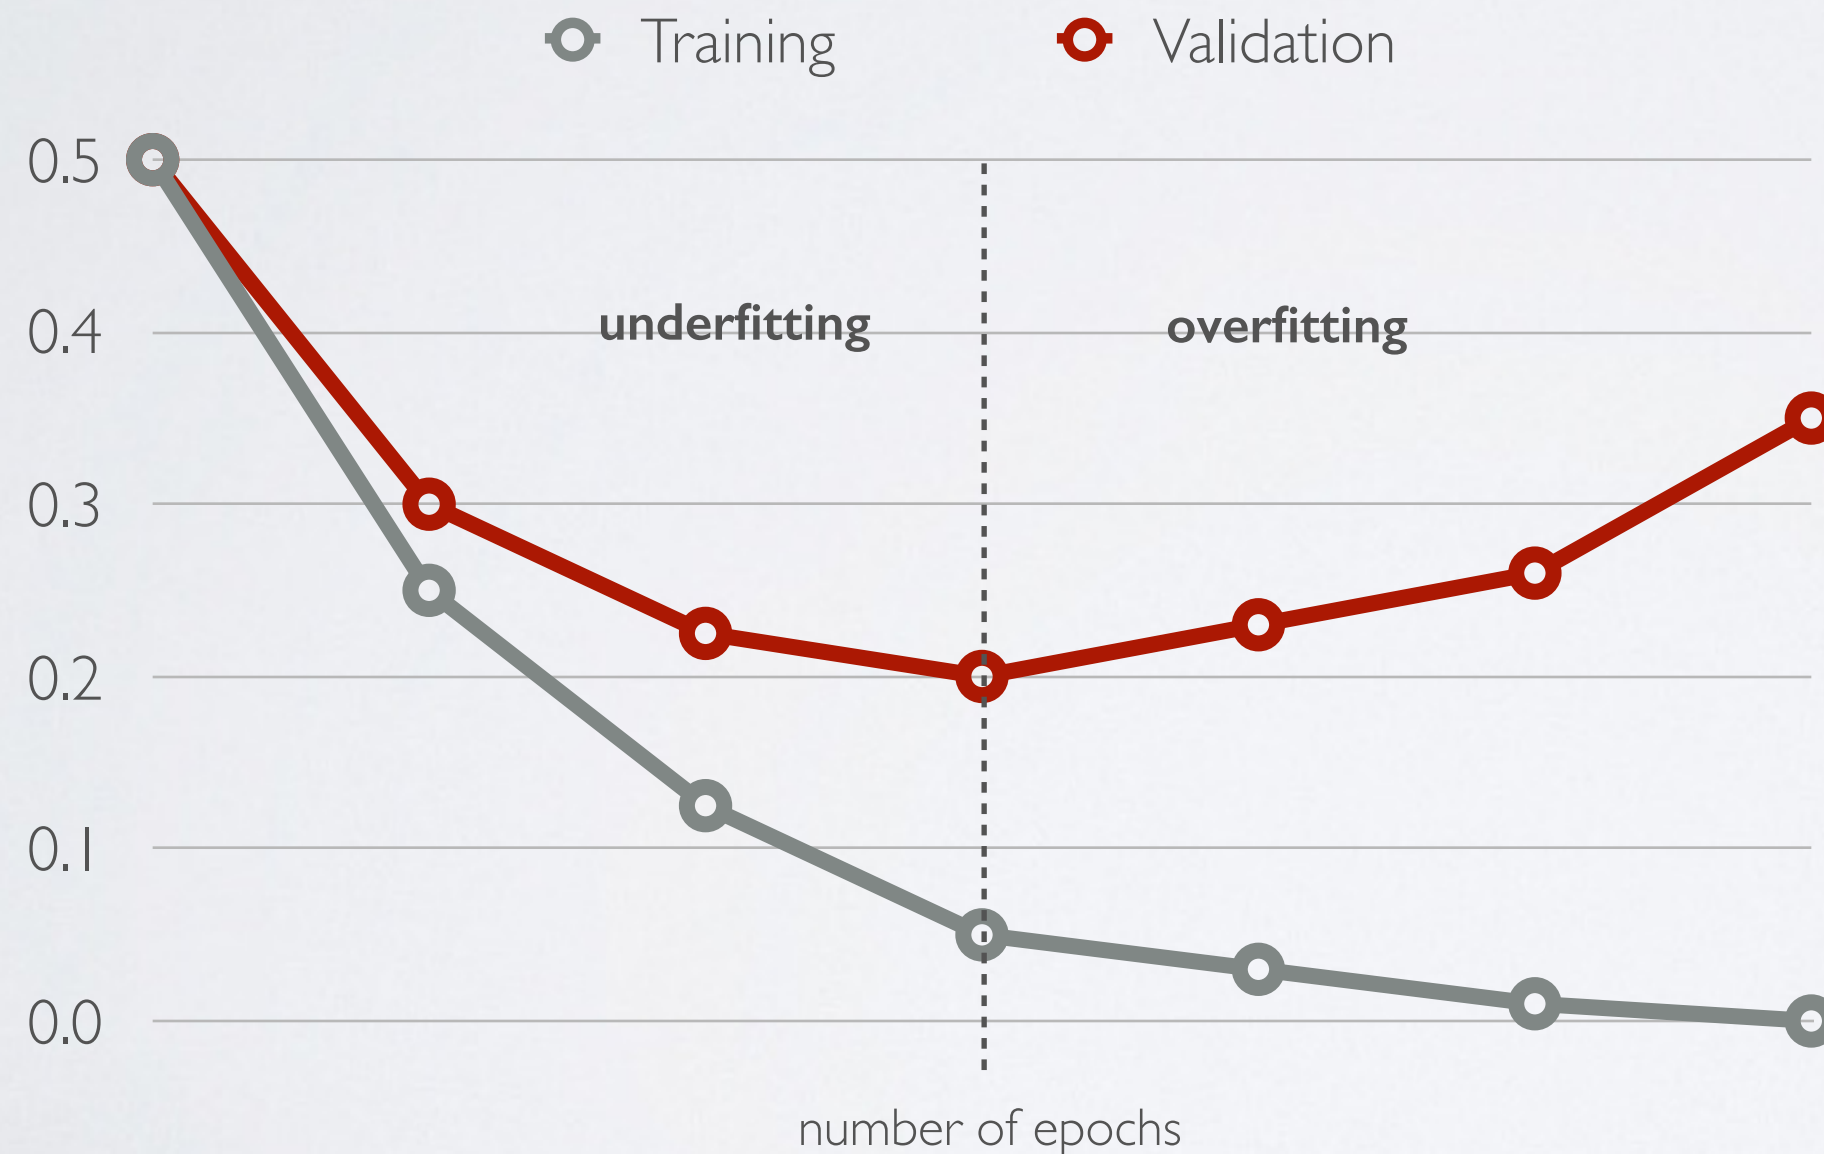




# KNOWING WHEN TO STOP


## Topics: early stopping

- To select the number of epochs, stop training when validation set error increases (with some look ahead)



# REGULARIZATION

## Topics: Early stopping in practice




---

**Algorithm 1** The early stopping meta-algorithm for determining the best amount of time to train. This meta-algorithm is a general strategy that works well with a variety of training algorithms and ways of quantifying error on the validation set.

---

Let  $n$  be the number of steps between evaluations.

Let  $p$  be the “patience,” the number of times to observe worsening validation set error before giving up.

Let  $\theta_o$  be the initial parameters.

$\theta \leftarrow \theta_o$

$i \leftarrow 0$

$j \leftarrow 0$

$v \leftarrow \infty$

$\theta^* \leftarrow \theta$

$i^* \leftarrow i$

**while**  $j < p$  **do**

    Update  $\theta$  by running the training algorithm for  $n$  steps.

$i \leftarrow i + n$

$v' \leftarrow \text{ValidationSetError}(\theta)$

**if**  $v' < v$  **then**

$j \leftarrow 0$

$\theta^* \leftarrow \theta$

$i^* \leftarrow i$

$v \leftarrow v'$

**else**

$j \leftarrow j + 1$

**end if**

**end while**

Best parameters are  $\theta^*$ , best number of training steps is  $i^*$

---

# REGULARIZATION

**Topics:** Early stopping with retraining

- Sometimes you really don't want to “waste” the validation set by not training on it.
- There are two basic strategies for retraining with the validation data.
  1. Retrain with train+valid for the same number of (updates / epochs) as determined by initial early stopping.
  2. Continue training w/ train+valid until the loss on valid = early-stopped loss on train. Not guaranteed to stop.

---

**Algorithm 1** A meta-algorithm for using early stopping to determine how long to train, then retraining on all the data.

---

Let  $\mathbf{X}^{(\text{train})}$  and  $\mathbf{y}^{(\text{train})}$  be the training set  
 Split  $\mathbf{X}^{(\text{train})}$  and  $\mathbf{y}^{(\text{train})}$  into  $\mathbf{X}^{(\text{subtrain})}$ ,  $\mathbf{y}^{(\text{subtrain})}$ ,  $\mathbf{X}^{(\text{valid})}$ ,  $\mathbf{y}^{(\text{valid})}$   
 Run early stopping starting from random  $\theta$  using  $\mathbf{X}^{(\text{subtrain})}$  and  $\mathbf{y}^{(\text{subtrain})}$  for training data and  $\mathbf{X}^{(\text{valid})}$  and  $\mathbf{y}^{(\text{valid})}$  for validation data. This returns  $i^*$ , the optimal number of steps.  
 Set  $\theta$  to random values again  
 Train on  $\mathbf{X}^{(\text{train})}$  and  $\mathbf{y}^{(\text{train})}$  for  $i^*$  steps.

---



---

**Algorithm 2** A meta-algorithm for using early stopping to determining at what objective value we start to overfit, then continuing training.

---

Let  $\mathbf{X}^{(\text{train})}$  and  $\mathbf{y}^{(\text{train})}$  be the training set  
 Split  $\mathbf{X}^{(\text{train})}$  and  $\mathbf{y}^{(\text{train})}$  into  $\mathbf{X}^{(\text{subtrain})}$ ,  $\mathbf{y}^{(\text{subtrain})}$ ,  $\mathbf{X}^{(\text{valid})}$ ,  $\mathbf{y}^{(\text{valid})}$   
 Run early stopping (Alg. ??) starting from random  $\theta$  using  $\mathbf{X}^{(\text{subtrain})}$  and  $\mathbf{y}^{(\text{subtrain})}$  for training data and  $\mathbf{X}^{(\text{valid})}$  and  $\mathbf{y}^{(\text{valid})}$  for validation data. This updates  $\theta$   
 $\epsilon \leftarrow J(\theta, \mathbf{X}^{(\text{subtrain})}, \mathbf{y}^{(\text{subtrain})})$   
**while**  $J(\theta, \mathbf{X}^{(\text{valid})}, \mathbf{y}^{(\text{valid})}) > \epsilon$  **do**  
   Train on  $\mathbf{X}^{(\text{train})}$  and  $\mathbf{y}^{(\text{train})}$  for  $n$  steps.  
**end while**

---

**Warning: these methods are dangerous!**



# REGULARIZATION

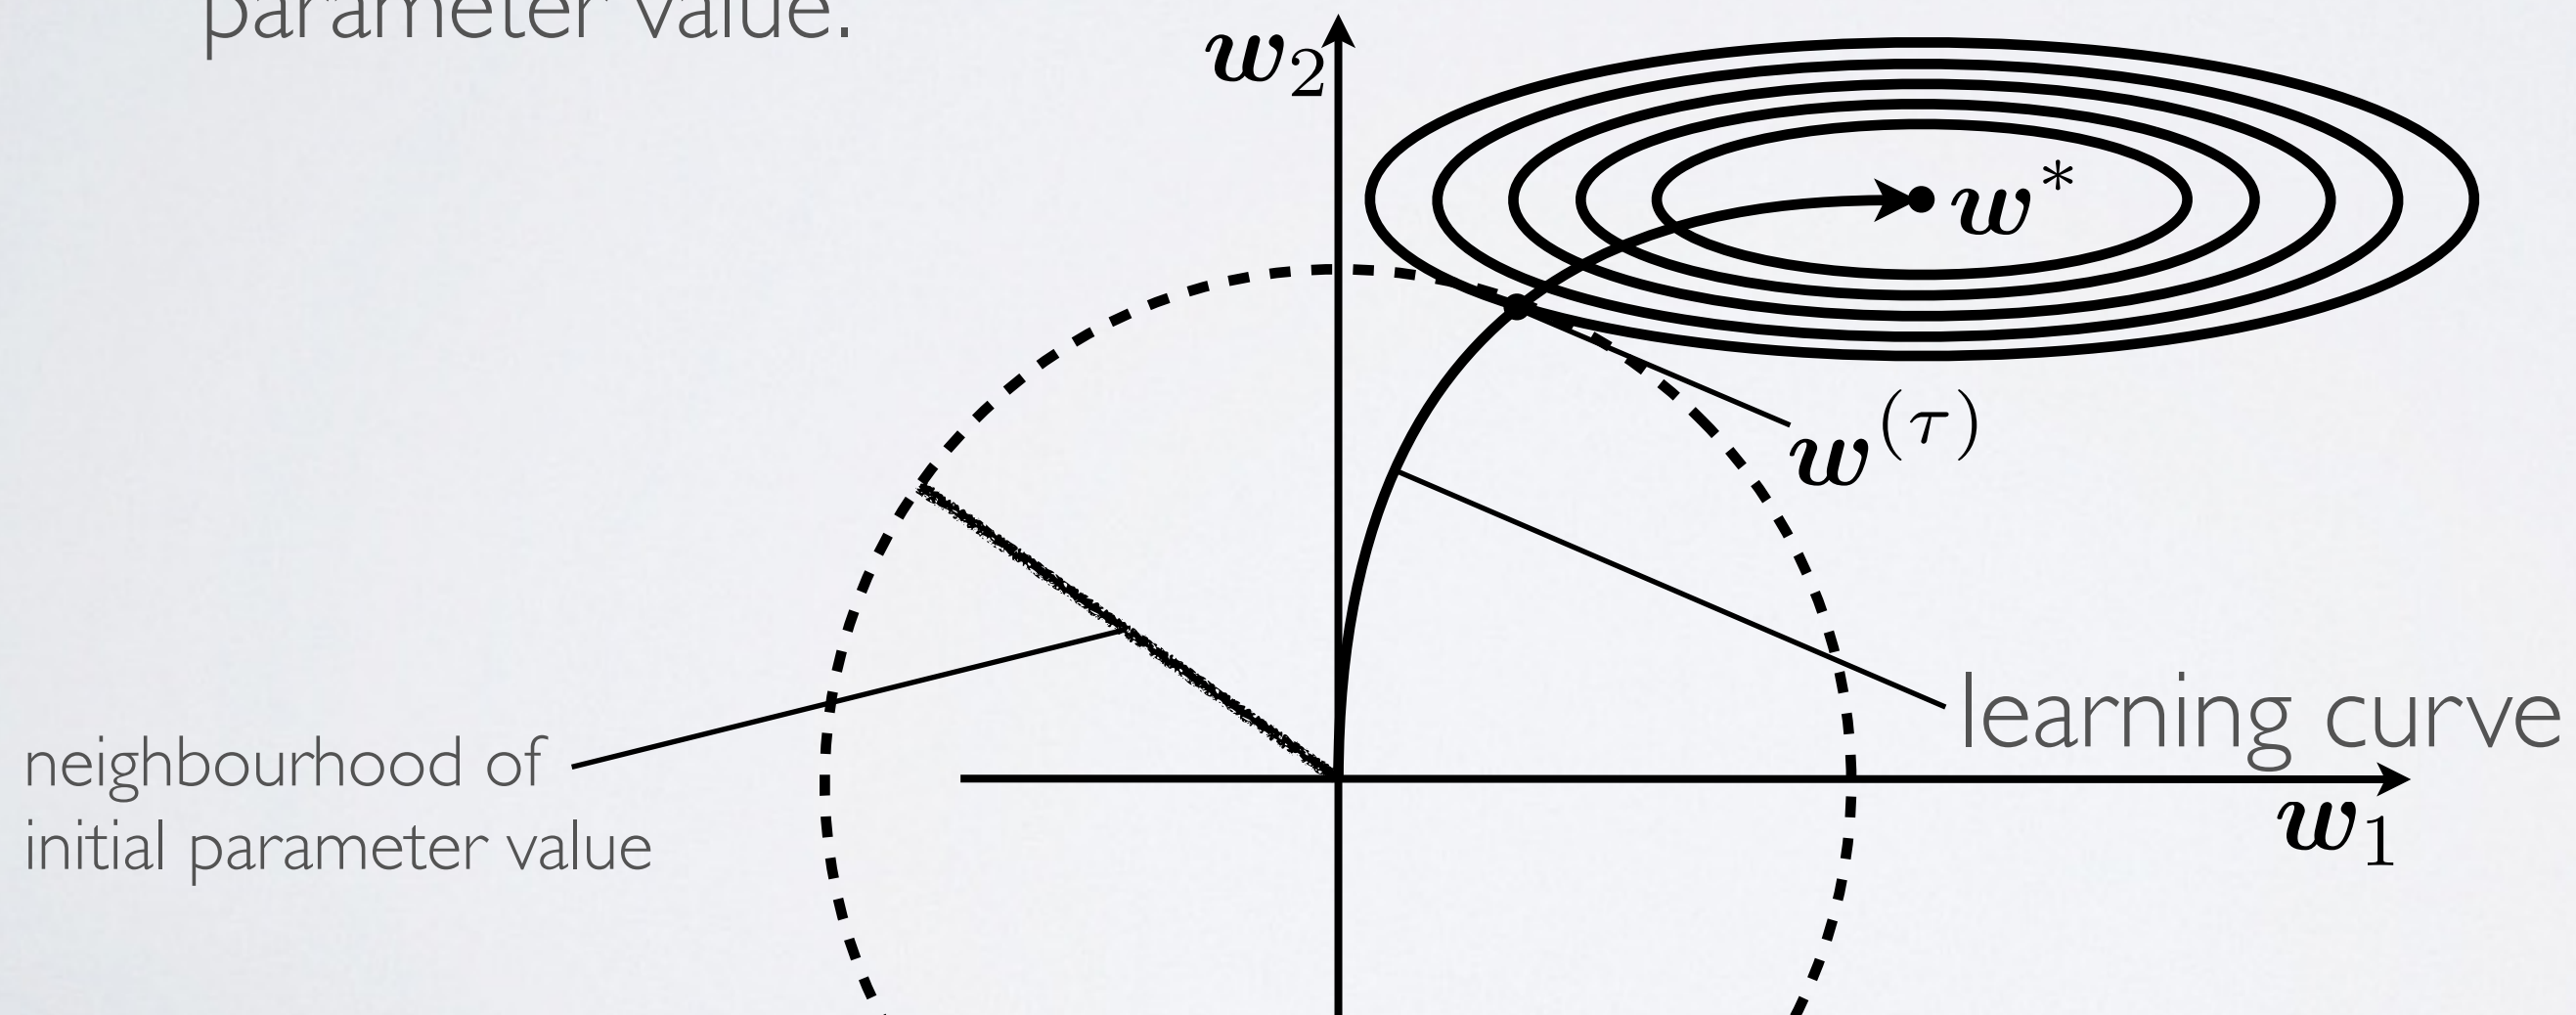
**Topics:** Early stopping with surrogate loss

- A useful property: can help to mitigate a mismatch between the surrogate loss and the underlying performance measure that we actually care about.
  - ▶ Example, 0-1 classification loss (derivative of zero almost everywhere). We therefore train with surrogates such as the log likelihood of correct class label.
  - ▶ However, 0-1 loss is inexpensive to compute, so it can easily be used as an early stopping criterion.
  - ▶ Often the 0-1 loss decreases long after the log likelihood has begun to worsen on the validation set.

# REGULARIZATION

**Topics:** How early stopping acts as a regularizer.

- What is the actual mechanism by which early stopping regularizes the model?
  - ▶ Early stopping has the effect of restricting the optimization procedure to a relatively small volume of parameter space in the neighbourhood of the initial parameter value.

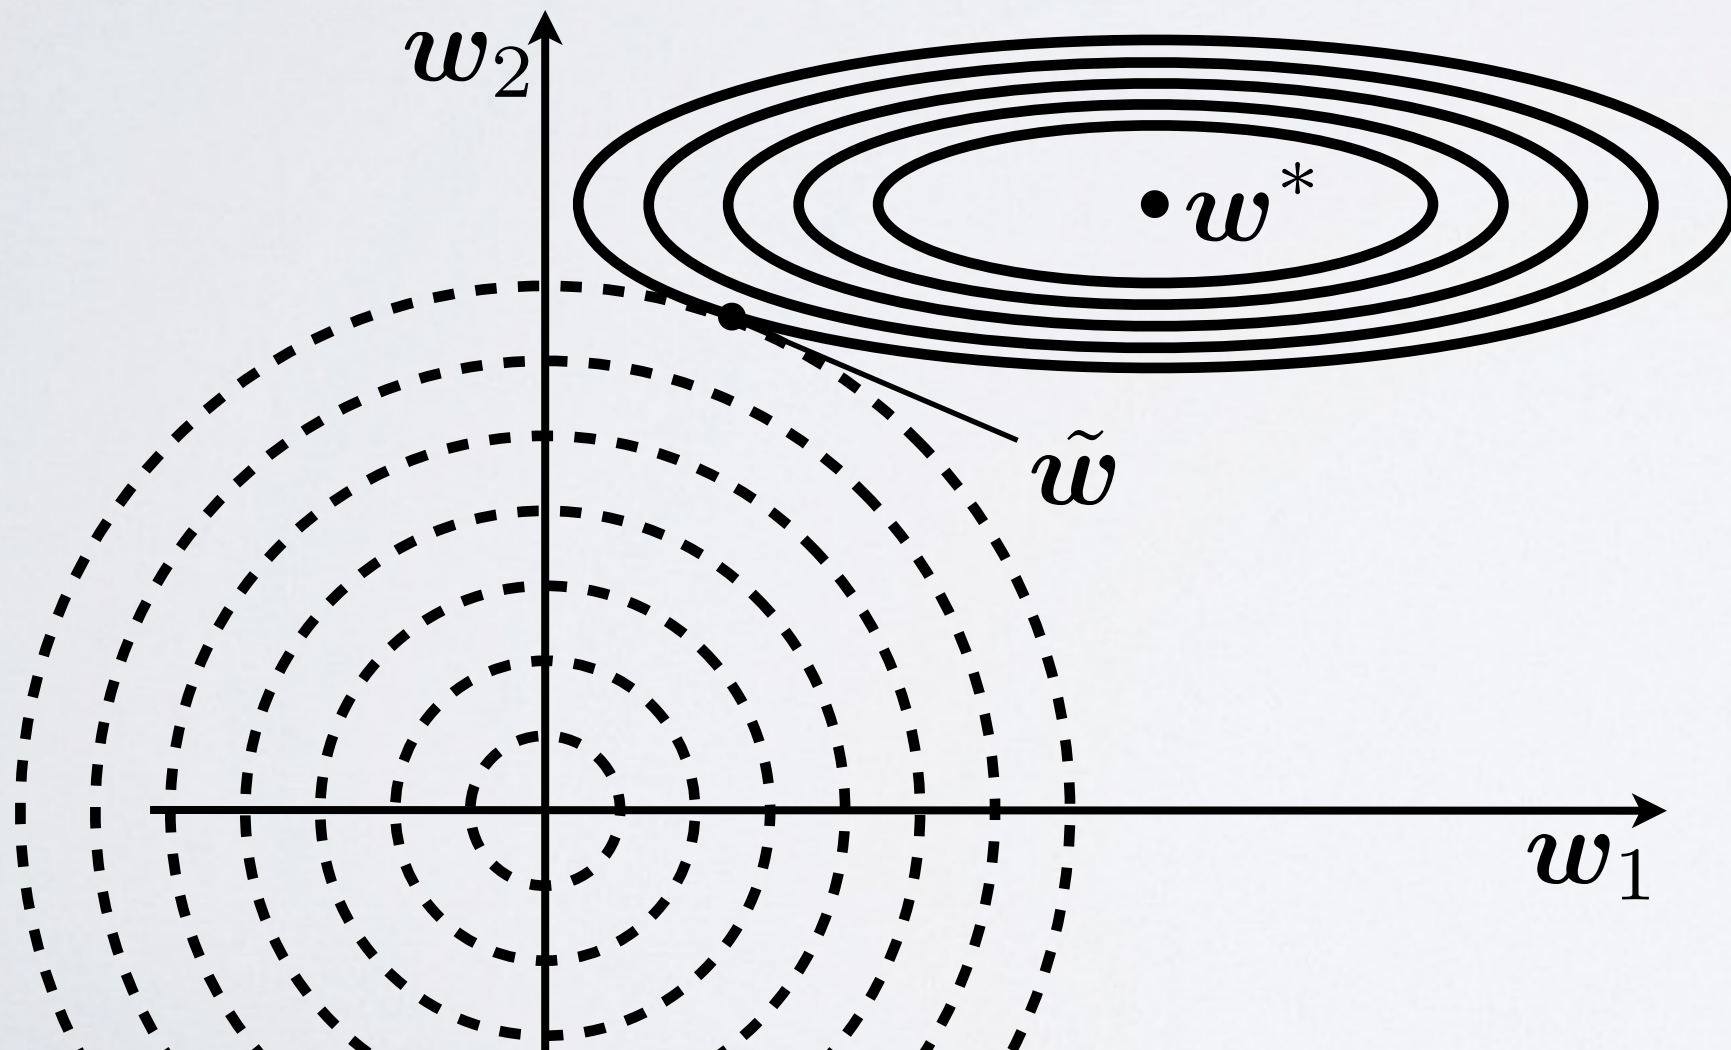


# REGULARIZATION

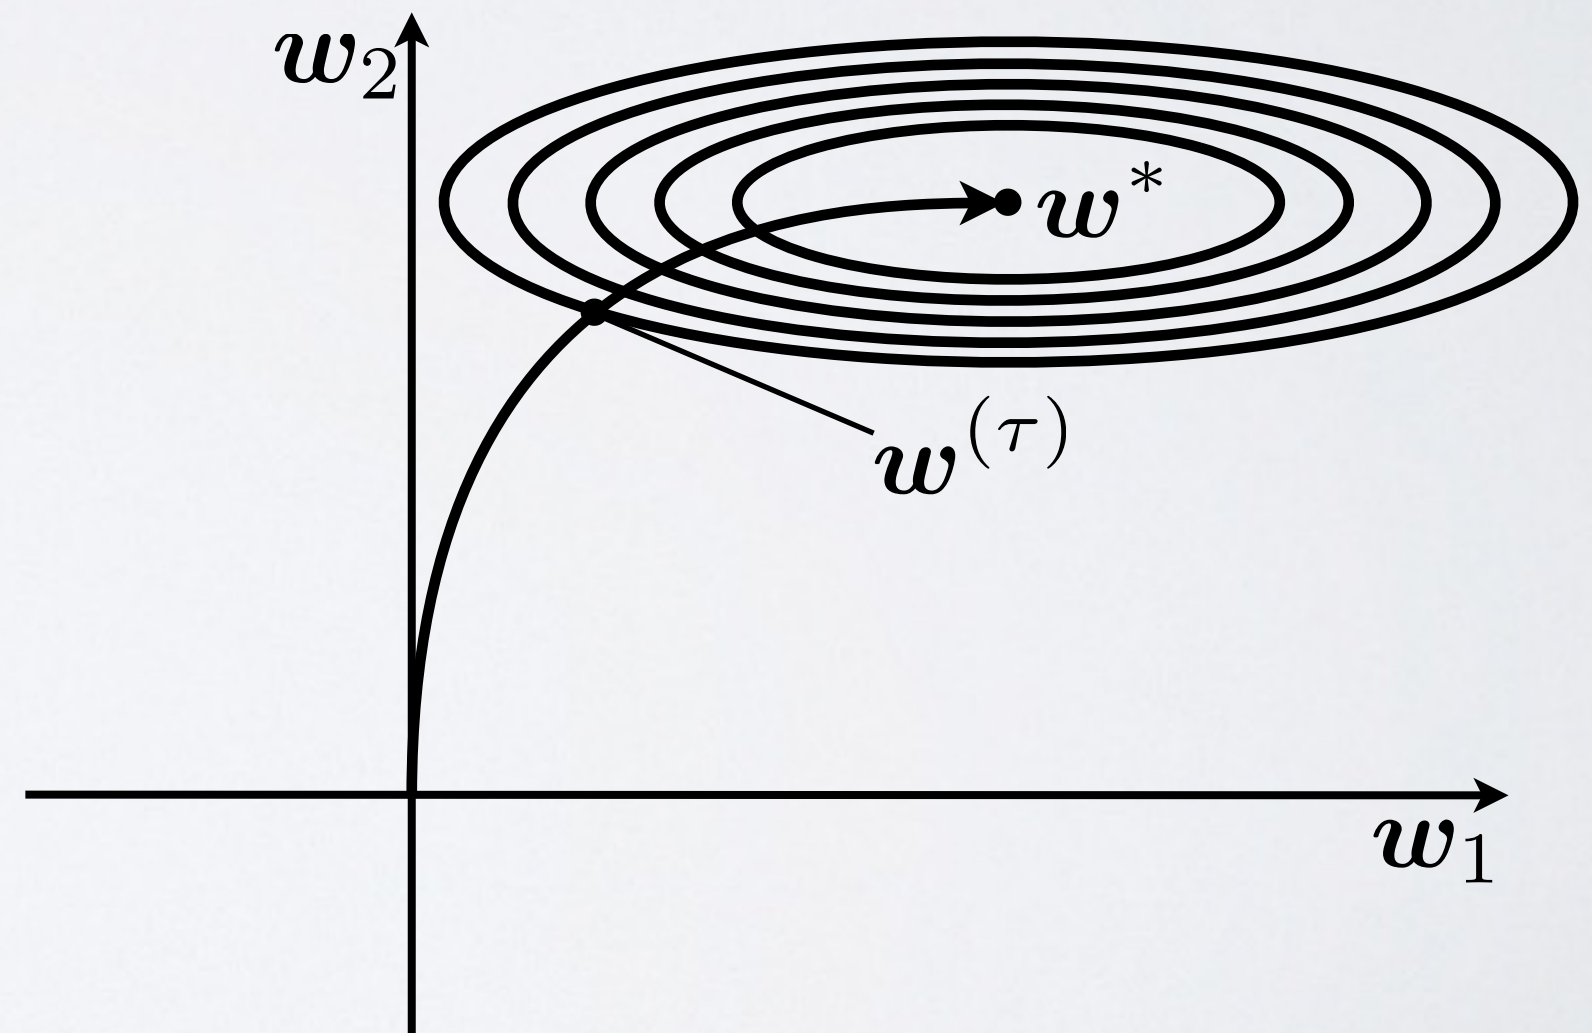
**Topics:** How early stopping acts as a regularizer.

- Assuming a simple linear model with a quadratic error function and simple gradient descent -- early stopping is equivalent to L2 regularization.

## L2 regularization



## Early Stopping





# REGULARIZATION

**Topics:** Early stopping equivalence to L2 regularization, **mathematical details**.

- Consider a quadratic approximation to the loss function in the neighbourhood of the empirically optimal value of the weights  $\boldsymbol{w}^*$

$$\hat{J}(\boldsymbol{\theta}) = J(\boldsymbol{w}^*) + \frac{1}{2}(\boldsymbol{w} - \boldsymbol{w}^*)^\top \boldsymbol{H}(\boldsymbol{w} - \boldsymbol{w}^*)$$

↙ Taking gradient

$$\nabla_{\boldsymbol{w}} \hat{J}(\boldsymbol{w}) = \boldsymbol{H}(\boldsymbol{w} - \boldsymbol{w}^*)$$

# REGULARIZATION

**Topics:** Early stopping equivalence to L2 regularization, **mathematical details.**

$$\nabla_{\boldsymbol{w}} \hat{J}(\boldsymbol{w}) = \boldsymbol{H}(\boldsymbol{w} - \boldsymbol{w}^*)$$

- Let us consider initial parameter vector chosen at the origin,
- We will consider updating the parameters via gradient descent:

$$\begin{aligned}\boldsymbol{w}^{(\tau)} &= \boldsymbol{w}^{(\tau-1)} - \eta \nabla_{\boldsymbol{w}} J(\boldsymbol{w}^{(\tau-1)}) \\ &= \boldsymbol{w}^{(\tau-1)} - \eta \boldsymbol{H}(\boldsymbol{w}^{(\tau-1)} - \boldsymbol{w}^*)\end{aligned}$$

$$\boldsymbol{w}^{(\tau)} - \boldsymbol{w}^* = (\boldsymbol{I} - \eta \boldsymbol{H})(\boldsymbol{w}^{(\tau-1)} - \boldsymbol{w}^*)$$

# REGULARIZATION

**Topics:** Early stopping equivalence to L2 regularization, **mathematical details.**

$$\mathbf{w}^{(\tau)} - \mathbf{w}^* = (\mathbf{I} - \eta \mathbf{H})(\mathbf{w}^{(\tau-1)} - \mathbf{w}^*)$$

- $\mathbf{H}$  is real and symmetric, so we can decompose it into a diagonal matrix  $\mathbf{\Lambda}$  and an orthogonal basis of eigenvectors,  $\mathbf{Q}$ , such that:  $\mathbf{H} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^\top$

$$\mathbf{w}^{(\tau)} - \mathbf{w}^* = (\mathbf{I} - \eta \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^\top)(\mathbf{w}^{(\tau-1)} - \mathbf{w}^*)$$

$$\mathbf{Q}^\top(\mathbf{w}^{(\tau)} - \mathbf{w}^*) = (\mathbf{I} - \eta \mathbf{\Lambda})\mathbf{Q}^\top(\mathbf{w}^{(\tau-1)} - \mathbf{w}^*)$$

- Assuming that  $|1 - \eta \lambda_i| < 1$  and that  $\mathbf{w}^{(0)} = \mathbf{0}$ . After  $\tau$  steps:

$$\mathbf{Q}^\top \mathbf{w}^{(\tau)} = [\mathbf{I} - (\mathbf{I} - \eta \mathbf{\Lambda})^\tau] \mathbf{Q}^\top \mathbf{w}^*$$



# REGULARIZATION

**Topics:** Early stopping equivalence to L2 regularization, **mathematical details.**

$$\mathbf{Q}^\top \mathbf{w}^{(\tau)} = [\mathbf{I} - (\mathbf{I} - \eta \mathbf{\Lambda})^\tau] \mathbf{Q}^\top \mathbf{w}^*$$

- Recall the L2 regularized solution was:  $\tilde{\mathbf{w}} = \mathbf{Q}(\mathbf{\Lambda} + \alpha \mathbf{I})^{-1} \mathbf{\Lambda} \mathbf{Q}^\top \mathbf{w}^*$

$$\mathbf{Q}^\top \tilde{\mathbf{w}} = (\mathbf{\Lambda} + \alpha \mathbf{I})^{-1} \mathbf{\Lambda} \mathbf{Q}^\top \mathbf{w}^*$$

$$\mathbf{Q}^\top \tilde{\mathbf{w}} = [\mathbf{I} - (\mathbf{\Lambda} + \alpha \mathbf{I})^{-1} \alpha] \mathbf{Q}^\top \mathbf{w}^*$$

- These are **equivalent** when  $(\mathbf{I} - \eta \mathbf{\Lambda})^\tau = (\mathbf{\Lambda} + \alpha \mathbf{I})^{-1} \alpha$

$$\tau \log(\mathbf{I} - \eta \mathbf{\Lambda}) = -\log(\mathbf{I} + \mathbf{\Lambda}/\alpha)$$

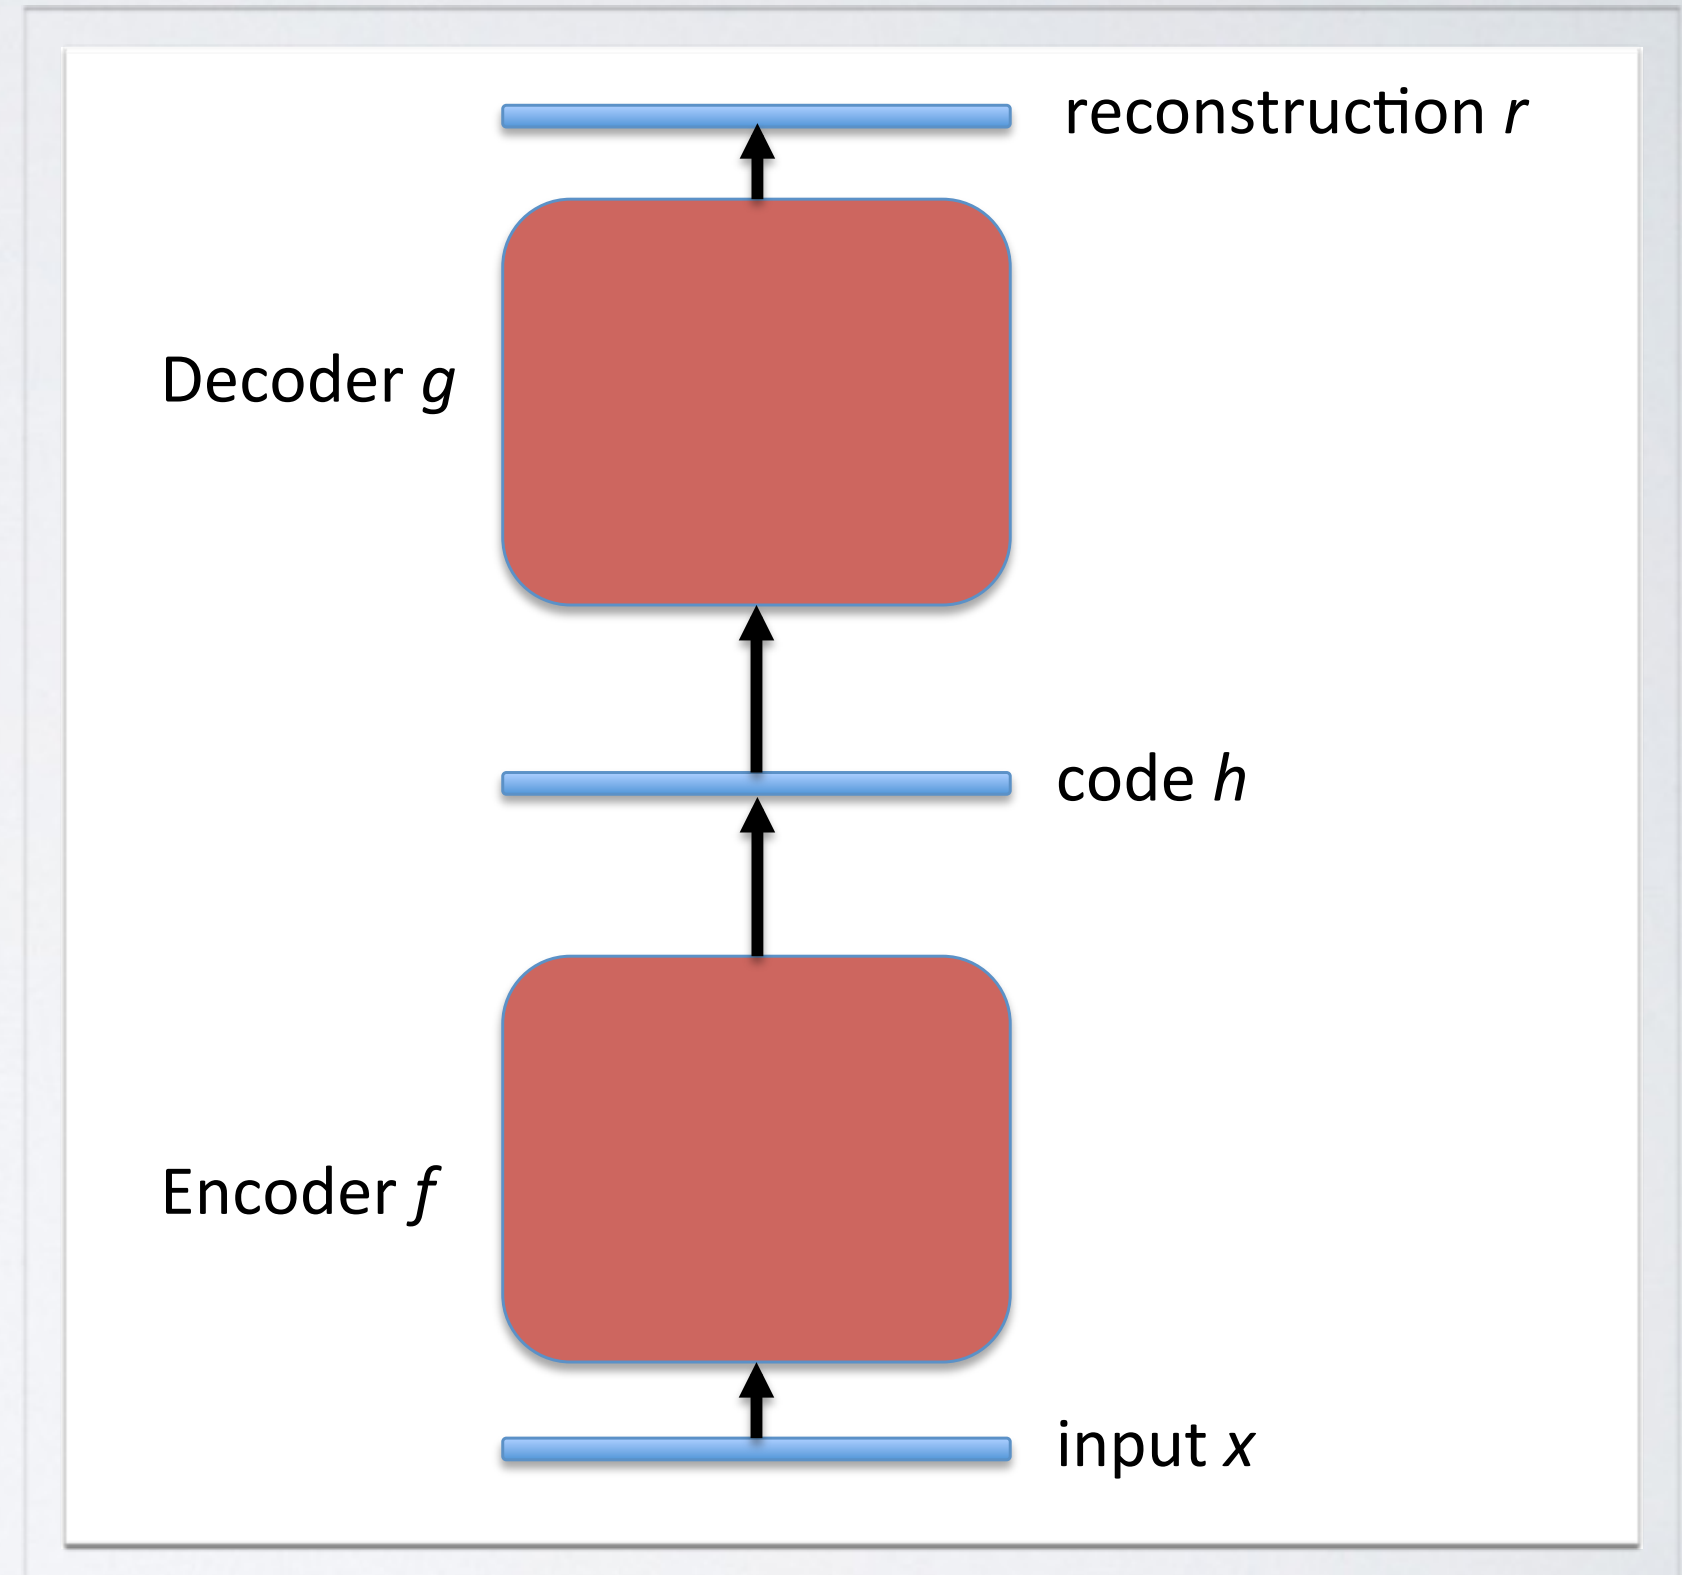
(by Taylor series expansion)  $\tau \approx 1/\eta\alpha$  for small  $\lambda_i \ \forall i$

Unsupervised learning as a  
regularization strategy

# REGULARIZATION

**Topics:** Unsupervised pretraining.

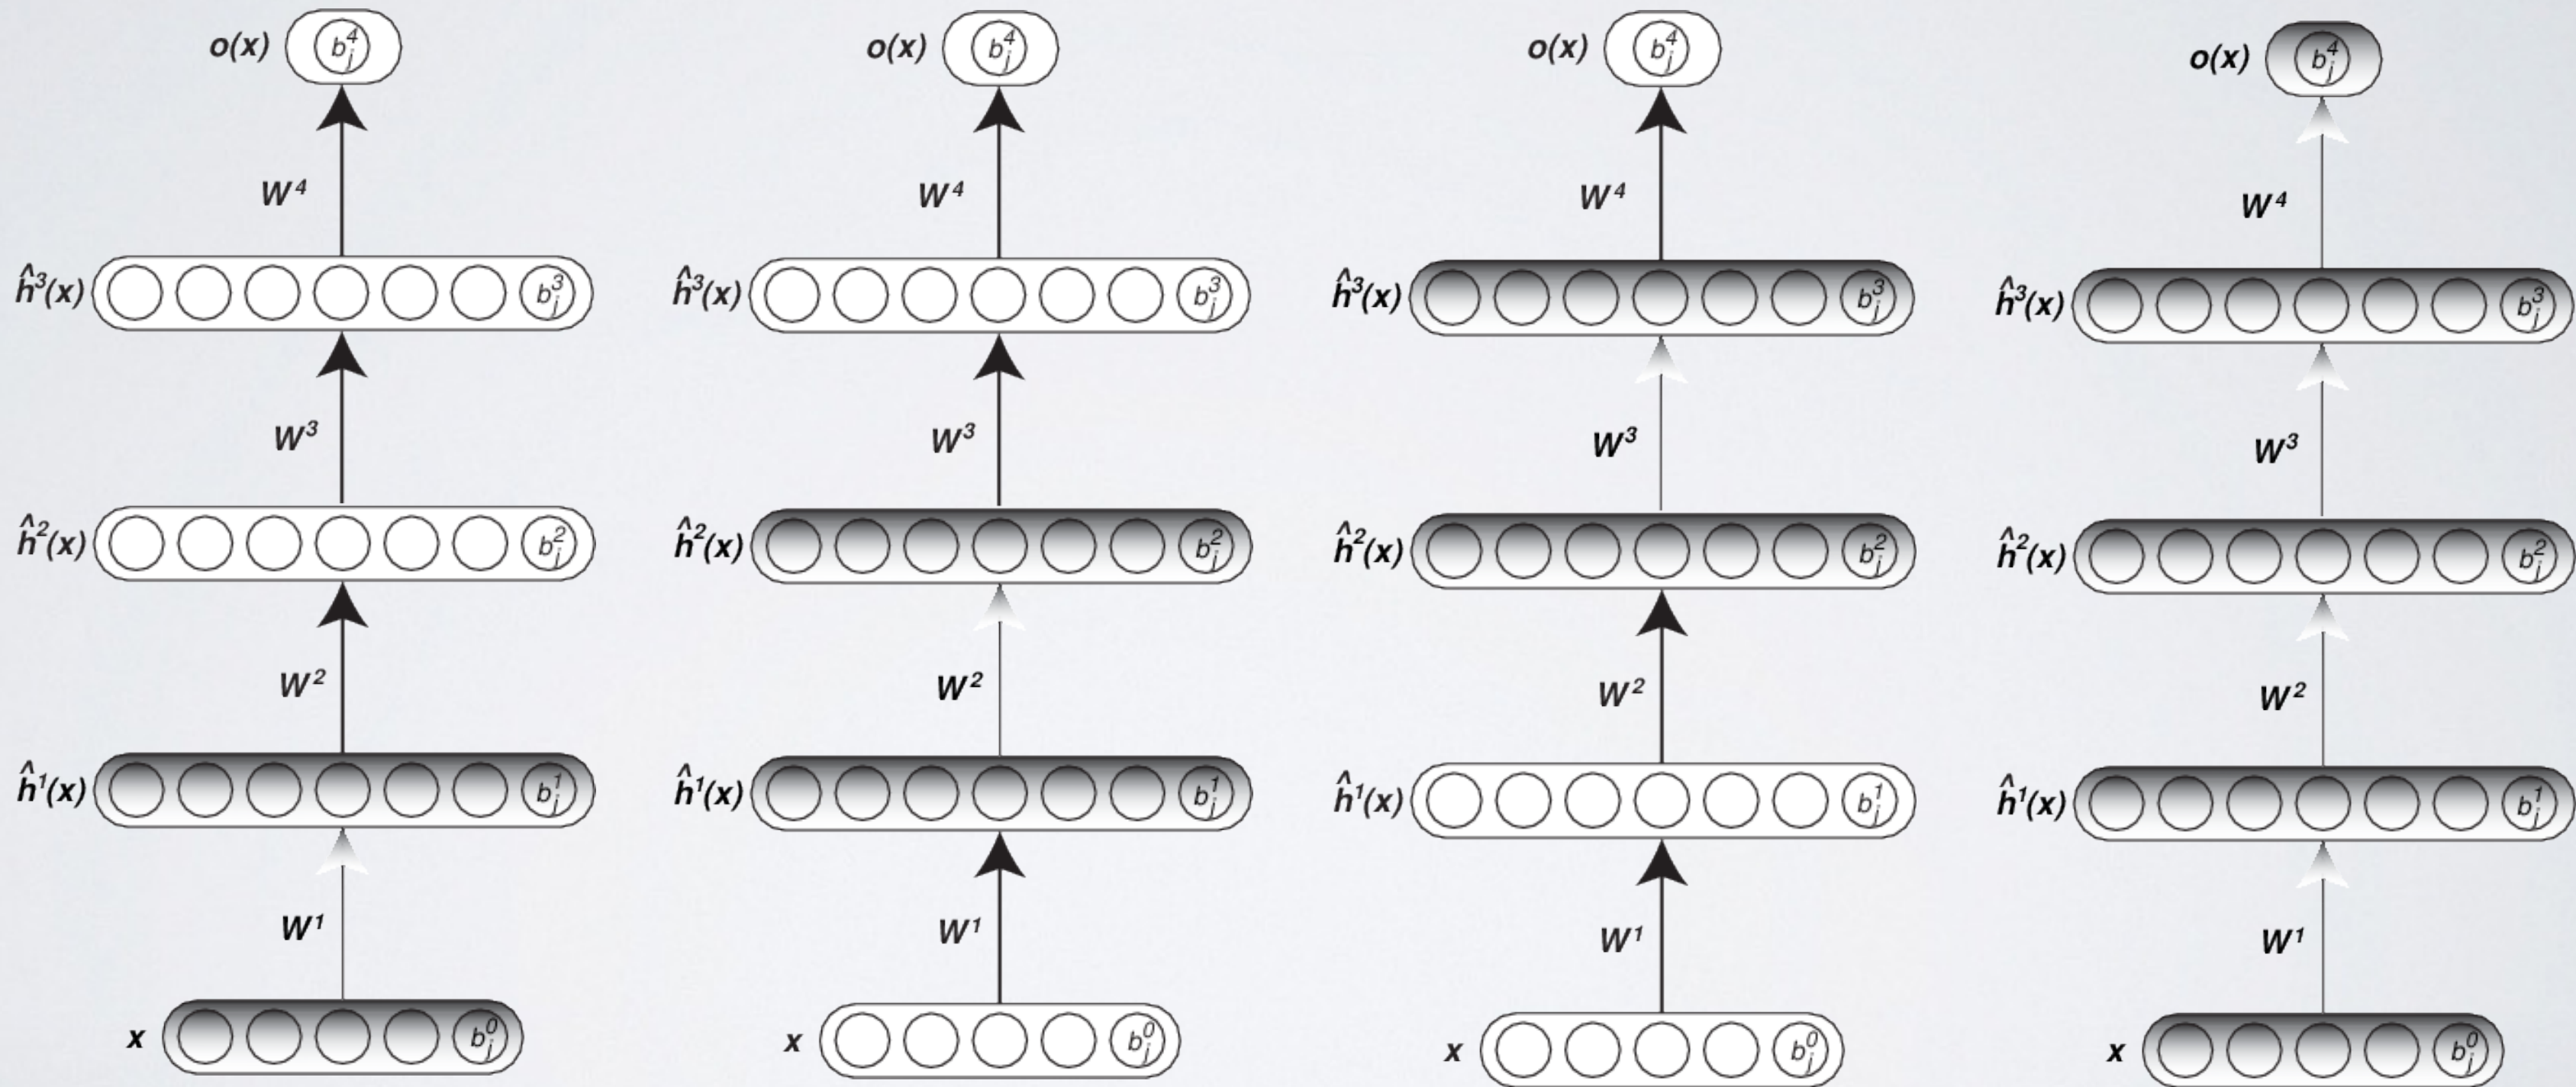
- Idea: pretrain your discriminative model parameters as an autoencoder.
- **Autoencoders** are featured prominently in the deep learning literature
- Goal: learn an encoder ( $f$ ) and decoder ( $g$ ) to minimize reconstruction error.
- Often, an additional penalty term is used to give the code ( $h$ ) desirable characteristics (we will see this later in the course)








# REGULARIZATION

**Topics:** Greedy layer-wise unsupervised pretraining.



# REGULARIZATION

- Topics:** Greedy layer-wise unsupervised pretraining as a regularization strategy:
- Training error / Test error profile matches that of a regularizer (Erhan et al. 2009).

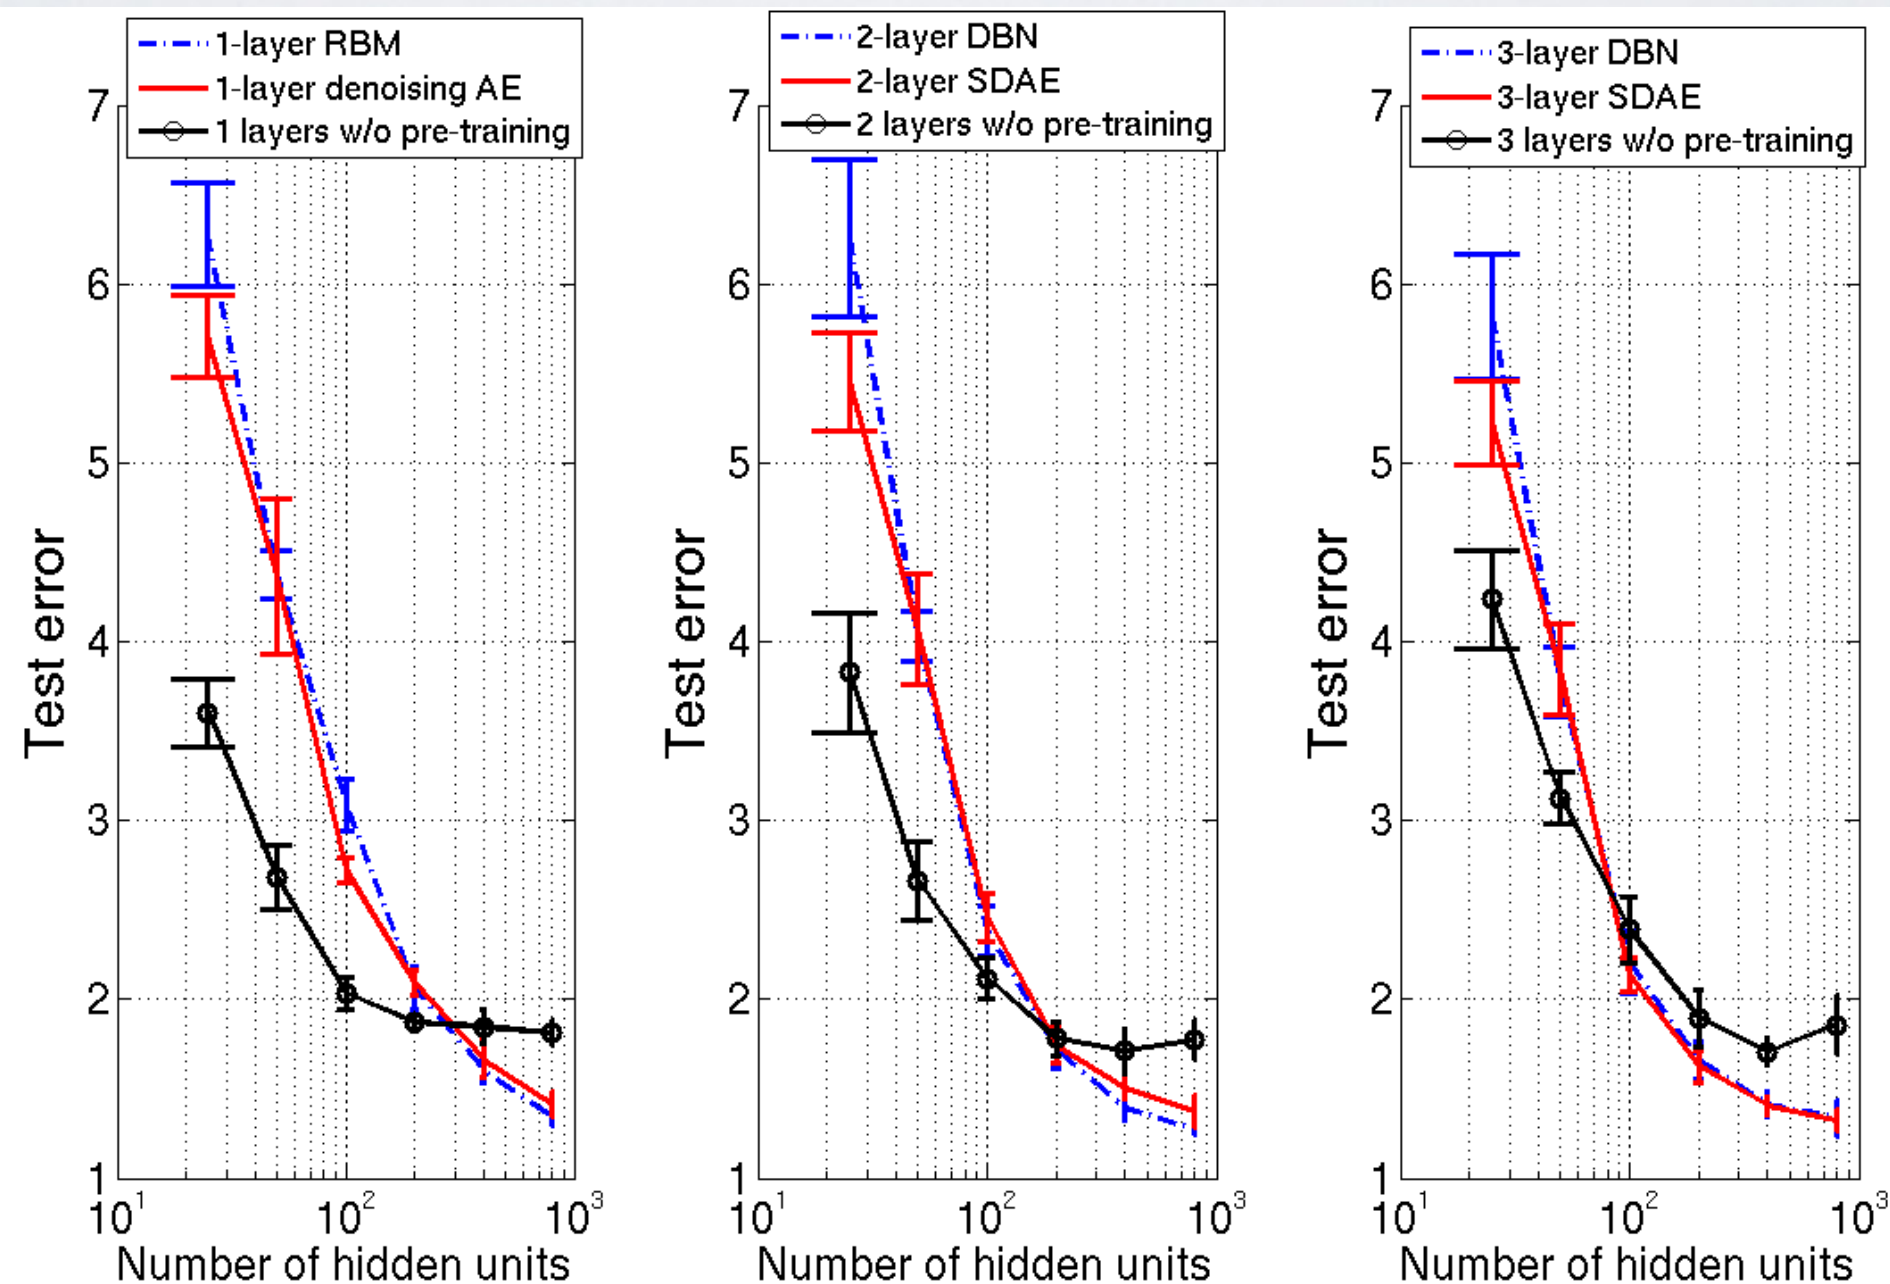




# REGULARIZATION

**Topics:** Greedy layer-wise unsupervised pretraining as a regularization strategy:

- Training error / Test error profile matches that of a regularizer (Erhan et al. 2009).

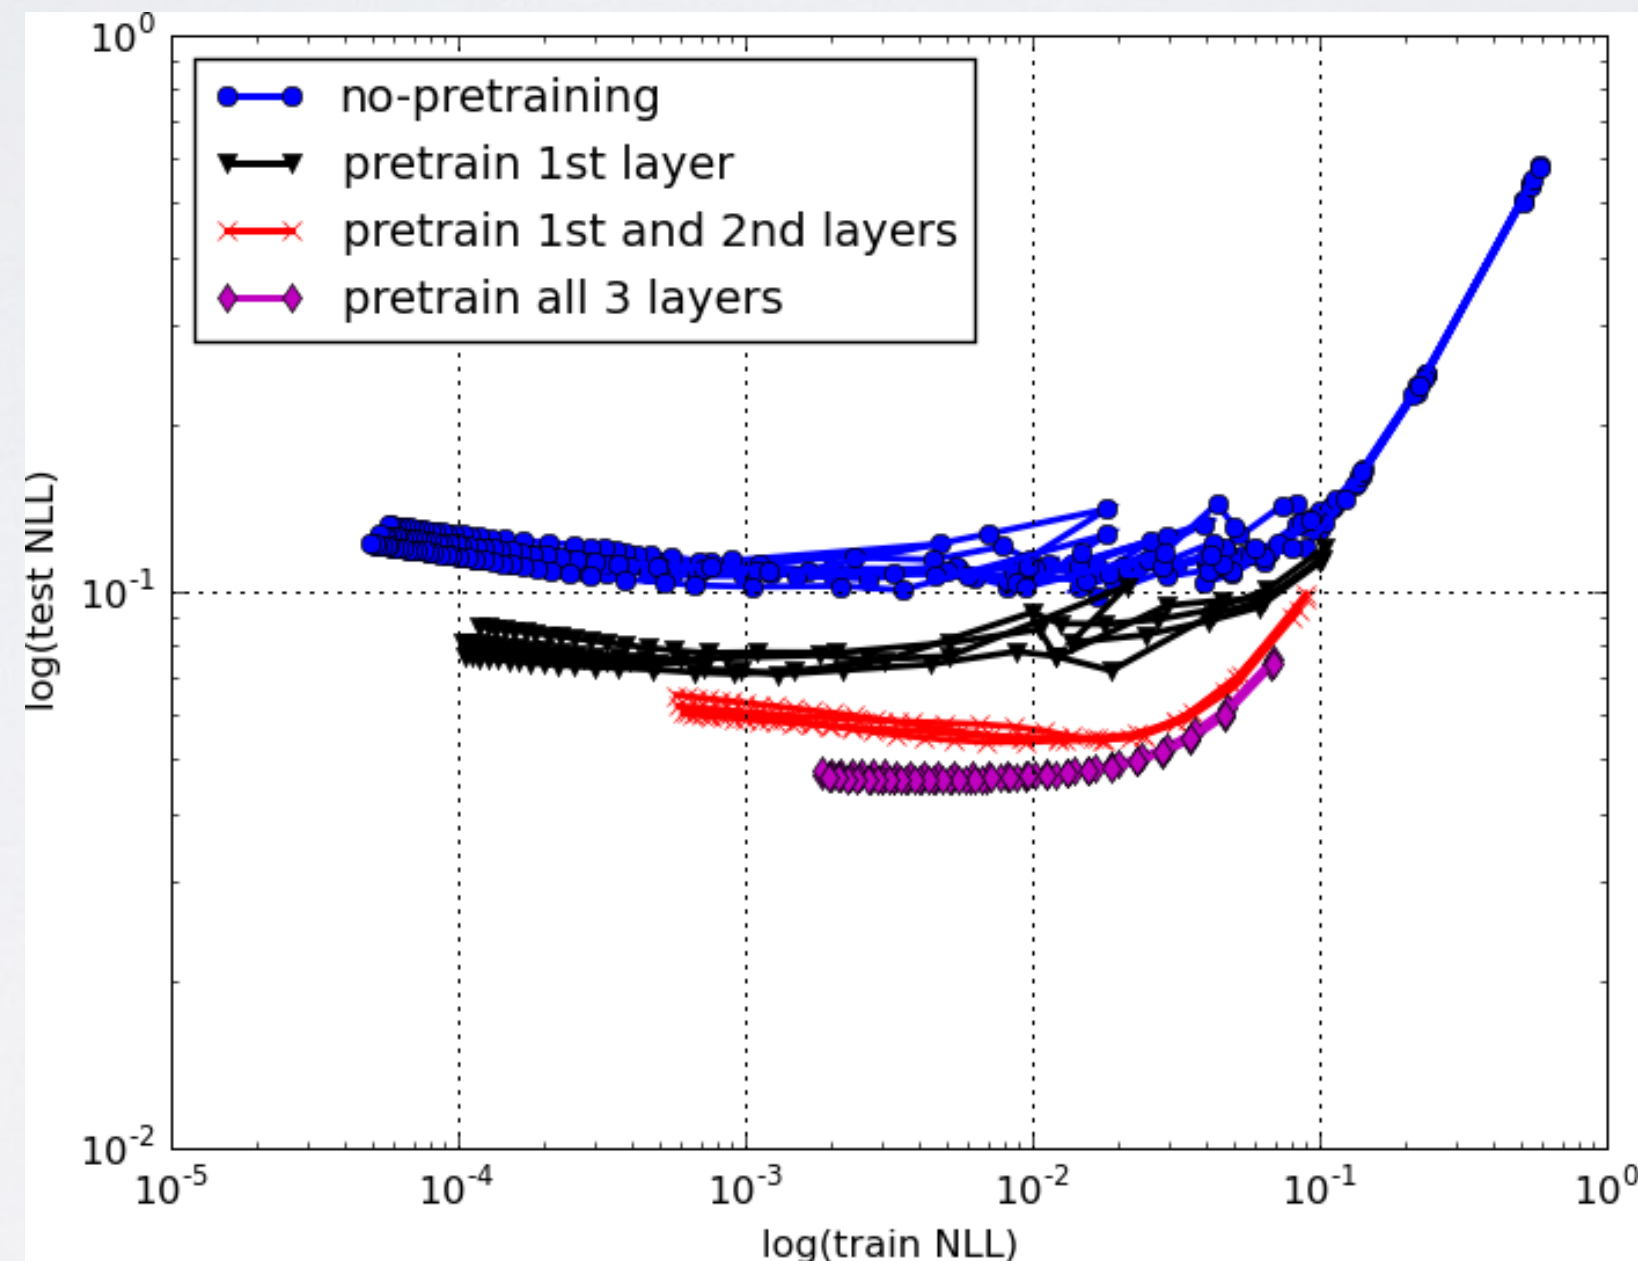




# REGULARIZATION

**Topics:** Greedy layer-wise unsupervised pretraining as a regularization strategy:

- Training error / Test error profile matches that of a regularizer (Erhan et al. 2009).



# REGULARIZATION

**Topics:** Multi-task learning / unsupervised learning.

- Same principle that applied to unsupervised learning applies to multi-task learning and transfer learning.
- Both are strategies to leverage other related tasks to **regularize** the parameters of the target task
- True even when there are multiple target tasks as in multi-task learning.
  - Each task regularizes the others.

